

STA261 Probability and Statistics II Winter 2019

Lecture Notes (Archived)

Yuchen Wang

March 15, 2025

Contents

1	Converge in distribution	3
2	Normal Distribution Theory	3
3	Expectation and Covariance	3
3.1	Expectation -Discrete case	3
3.2	Expectation - Continuous case	4
3.3	Variance, Covariance and Correlation	5
4	Independent Random Variables	6
5	Types of Inferences	7
6	Different Types of Estimation	7
6.1	Method of Moments Estimation	7
6.2	Maximum Likelihood Estimation	8
7	Sampling Distribution of an Estimator	10
8	Population Variance (σ^2)	11
9	Sampling distribution of S^2(under Normal Distribution	12
10	Some relationships among distributions	14
11	Difference between sample variance and variance of sample mean	14

12 Consistent Estimator	15
13 Efficient Estimator	16

1 Converge in distribution

2 Normal Distribution Theory

Theorem: Sum of independent normal random variables Suppose $X_i \sim N(\mu_i, \sigma_i^2)$ for $i = 1, 2, \dots, n$ and that they are independent random variables. Let $Y = (\sum_i a_i X_i) + b$ for some constants $\{a_i\}$ and b . Then

$$Y \sim N((\sum_i a_i \mu_i) + b, \sum_i a_i^2 \sigma_i^2)$$

Corollary: The distribution of the sample mean of normal random variables Suppose $X_i \sim N(\mu, \sigma^2)$ for $i = 1, 2, \dots, n$ and that they are independent random variables, If $\bar{X} = (X_1 + \dots + X_n)/n$, then $\bar{X} \sim N(\mu, \sigma^2/n)$

Theorem: The covariance of sums of normal random variables Suppose $X_i \sim N(\mu_i, \sigma_i^2)$ for $i = 1, 2, \dots, n$ and also that the $\{X_i\}$ are independent. Let $U = \sum_{i=1}^n a_i X_i$ and $V = \sum_{i=1}^n b_i X_i$ for some constants $\{a_i\}$ and $\{b_i\}$. Then $Cov(U, V) = \sum_i a_i b_i \sigma_i^2$. Furthermore, $Cov(U, V) = 0$ if and only if U and V are independent.

3 Expectation and Covariance

3.1 Expectation -Discrete case

Definition of expectation Let X be a discrete random variable, taking on discrete values x_1, x_2, \dots , with $p_i = P(X = x_i)$. Then the *expected value* (or *mean* or *mean value*) of X , written $E(X)$ (or μ_x), is defined by

$$E(X) = \sum_i x_i p_i$$

Theorem: expectation involving nested functions

1. Let X be a discrete random variable, and let $g : \mathbb{R} \rightarrow \mathbb{R}$ be some function such that the expectation of the random variable $g(X)$ exists. Then

$$E(g(X)) = \sum_x g(x) P(X = x)$$

2. Let X and Y be discrete random variables, and let $h : \mathbb{R}^2 \rightarrow \mathbb{R}$ be some function such that the expectation of the random variable $h(X, Y)$ exists. Then

$$E(h(X, Y)) = \sum_{x,y} h(x, y) P(X = x, Y = y)$$

Theorem: Linearity of expected values Let X and Y be discrete random variables, let a and b be real numbers, and put $Z = aX + bY$. Then

$$E(Z) = aE(X) + bE(Y)$$

Theorem: Expectation of product of independent r.v Let X and Y be discrete random variables that are independent. Then

$$E(XY) = E(X)E(Y)$$

Monotonicity Let X and Y be discrete random variables, and suppose that $X \leq Y$ (Remember that this means $X(s) \leq Y(s)$ for all $s \in S$) Then $E(X) \leq E(Y)$.

3.2 Expectation - Continuous case

Definition of expectation Let X be an absolutely continuous random variable, with density function f_X . Then the *expected value* of X is given by

$$E(x) = \int_{-\infty}^{\infty} xf_X(x)dx$$

Theorem: expectation involving nested functions

1. Let X be a an absolutely continuous random variable with density function f_X , and let $g : \mathbb{R} \rightarrow \mathbb{R}$ be some function such that the expectation of the random variable $g(X)$ exists. Then

$$E[g(x)] = \int_{-\infty}^{\infty} g(x)f_X(x)dx$$

2. Let X and Y be discrete random variables, and let $h : \mathbb{R}^2 \rightarrow \mathbb{R}$ be some function such that the expectation of the random variable $h(X, Y)$ exists. Then

$$E(h(X, Y)) = \int_{-\infty}^{\infty} h(x, y)f_{X,Y}(x, y)dxdy$$

Theorem: Linearity of expected values Let X and Y be jointly absolutely continuous random variables, let a and b be real numbers. Then

$$E(aX + bY) = aE(X) + bE(Y)$$

Monotonicity Let X and Y be jointly continuous random variables, and suppose that $X \leq Y$ (Remember that this means $X(s) \leq Y(s)$ for all $s \in S$) Then $E(X) \leq E(Y)$.

3.3 Variance, Covariance and Correlation

Definition of variance The *variance* of a random variable X is the quantity

$$\sigma_x^2 = \text{Var}(X) = E((X - \mu_X)^2)$$

where σ_X is the *standard deviation* of X .

Theorem Let X be any r.v. with $\mu_X = E(X)$ and variance $\text{Var}(X)$. Then the following hold true:

1. $\text{Var}(X) \geq 0$
2. If a and b are real numbers, $\text{Var}(aX + b) = a^2 \text{Var}(X)$
3. $\text{Var}(X) = E(X^2) - (\mu_X)^2 = E(X^2) - E(X)^2$
4. $\text{Var}(X) \leq E(X^2)$

Definition of covariance

$$\text{Cov}(X, Y) = E((X - \mu_X)(Y - \mu_Y))$$

Theorem: Linearity of covariance Let X , Y and z be three r.v.s. Let a and b be real numbers. Then

$$\text{Cov}(aX + bY, Z) = a\text{Cov}(X, Z) + b\text{Cov}(Y, Z)$$

Theorem Let X and Y be r.v.s. Then

$$\text{Cov}(X, Y) = E(XY) - E(X)E(Y)$$

Theorem If X and Y are independent, then

$$\text{Cov}(X, Y) = 0$$

Theorem

1. For any r.v.s X and Y ,

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$$

2. More generally, for any r.v.s X_1, \dots, X_n ,

$$\text{Var}(\sum_i X_i) = \sum_i \text{Var}(X_i) + 2\sum_{i < j} \text{Cov}(X_i, X_j)$$

Corollary

1. If X and Y are independent, then $\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y)$
2. If X_1, \dots, X_n are independent, then $\text{Var}(\sum_{i=1}^n X_i) = \sum_{i=1}^n \text{Var}(X_i)$

Definition The *correlation* of two r.v.s X and Y is given by

$$\text{Corr}(X, Y) = \frac{\text{Cov}(X, Y)}{\text{Sd}(X)\text{Sd}(Y)}$$

provided $\text{Var}(X) < \infty$ and $\text{Var}(Y) < \infty$

4 Independent Random Variables

Definition 1 Let X and Y be two continuous random variables. We say X and Y are independent if

$$f_{X,Y}(x, y) = f_X(x) \times f_Y(y)$$

$\forall x, y \in \mathbb{R}$

Lemma 1 X and Y are two continuous random variables. If X and Y are independent, then

$$E[g(X)h(Y)] = E(g(X)) \times E[h(Y)]$$

for any two functions $g()$ and $h()$.

proof:

$$\begin{aligned}
 E[g(X)h(Y)] &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} g(x)h(y)f_{X,Y}(x,y) dx dy \\
 &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} g(x)h(y)f_X(x)f_Y(y) dx dy \\
 &= \int_{-\infty}^{\infty} f_Y(y)h(y) \int_{-\infty}^{\infty} g(x)f_X(x) dx dy \\
 &= \int_{-\infty}^{\infty} f_Y(y)h(y)E[g(X)] dy \\
 &= E[g(X)] \int_{-\infty}^{\infty} f_Y(y)h(y) dy \\
 &= E[g(X)]E[h(Y)]
 \end{aligned}$$

■

5 Types of Inferences

Estimation:

1. Point estimation: Based on the sample observations, calculating a particular value as an estimate of the parameter θ
2. Interval estimation: Calculating a range of values that is likely to contain the parameter θ

Hypothesis testing Based on the sample, assess whether a hypothetical value θ_0 is a plausible value of the parameter θ or not.

6 Different Types of Estimation

6.1 Method of Moments Estimation

Let X_1, X_2, \dots, X_n are independently and identically distributed (i.i.d.) random variables.

Let the k^{th} population moment be

$$\mu_k = E[X^k]$$

k^{th} sample moment based on sample

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i^k$$

We use $\hat{\mu}_k$ as an estimator of μ_k

In other words, we use the sample moments as estimators of the population moments.

6.2 Maximum Likelihood Estimation

Definition of Likelihood Function Suppose X_1, X_2, \dots, X_n has a joint density or mass function $f(x_1, x_2, \dots, x_n | \theta)$

We observe sample, $X_1 = x_1, X_2 = x_2, \dots, X_n = x_n$

Given the sample, the likelihood function of θ , noted as $L(\theta | x_1, x_2, \dots, x_n)$, is defined as

$$L(\theta | x_1, x_2, \dots, x_n) = f(x_1, x_2, \dots, x_n | \theta)$$

Often written as $L(\theta)$, is a function of θ .

If X follows a discrete distribution, it gives the probability of observing the sample as a function of the parameter θ

If X_1, X_2, \dots, X_n are i.i.d. then their joint density is the product of marginal densities, $f_\theta(x)$

Hence, in i.i.d. case we write

$$L(\theta) = \prod_{i=1}^n f_\theta(x_i)$$

Comments

1. $L(\theta)$ is NOT a pdf or pmf of θ
2. Likelihood introduces a belief ordering on parameter space, Ω
3. For $\theta_1, \theta_2 \in \Omega$, we believe in θ_1 as the true value of θ over θ_2 whenever $L(\theta_1) > L(\theta_2)$
4. Which means, the data is more likely to come from f_{θ_1} than f_{θ_2}
5. The value $L(\theta)$ is very small for every value of θ
6. So often, we are interested in the likelihood ratios:

$$\frac{L(\theta_1)}{L(\theta_2)}$$

Maximum Likelihood Estimation

1. Let's say we are interested in a point estimate of θ
2. A sensible choice will be to pick $\hat{\theta}$ that maximizes $L(\theta)$
3. So $\hat{\theta}$ satisfies $L(\hat{\theta}) \geq L(\theta)$ for all $\theta \in \Omega$
4. $\hat{\theta}$ is called the maximum likelihood estimate (MLE) of θ

Computation of the MLE

1. Define, log-likelihood function, $l(\theta) = \ln L(\theta)$
2. $\ln(x)$ is a 1-1 increasing function of $x > 0 \implies L(\hat{\theta}) \geq L(\theta)$ for $\theta \in \Omega$ iff $l(\hat{\theta}) \geq l(\theta)$
3. In other words, if $L(\theta)$ is maximized at $\hat{\theta}$ then $l(\theta)$ will also be maximized at $\hat{\theta}$
4. Therefore,

$$l(\theta) = \ln(\prod_{i=1}^n f_{\theta}(x_i)) = \sum_{i=1}^n \ln f_{\theta}(x_i)$$

5. The obvious benefit: It's much easier to differentiate a sum than a product
6. Solve the equation, $\frac{\partial l(\theta)}{\partial \theta} = 0$ for θ
7. Say, $\hat{\theta}$ is the solution. But it's still not the MLE
8. Need to check whether or not

$$\frac{\partial^2 l(\theta)}{\partial \theta^2} \Big|_{\theta=\hat{\theta}} < 0$$

Properties of MLE

1. MLE is not unique
2. MLE may not exist
3. The likelihood may not always be differentiable.

7 Sampling Distribution of an Estimator

1. Recall: An Estimator (T) is a random variable (infinite number of sample means)
2. If we repeat the sampling procedure and keep calculating T for each set of sample and finally draw a density histogram based on the T values we get the sampling distribution of T
3. **Standard error:** Standard deviation of an estimator is called the standard error (SE)

Definition of Mean Squared Error Let $\psi(\theta)$ be any real valued function of θ , suppose T is an estimator of $\psi(\theta)$

$$MSE_{\theta}(T) = E_{\theta}[(T - \psi(\theta))^2]$$

Corollary

$$MSE_{\theta}(T) = Var_{\theta}(T) + (E_{\theta}(T) - \psi(\theta))^2$$

proof:

$$\begin{aligned} MST(T) &= E[(T - \psi(\theta))^2] \\ &= E[(T - E(T) + E(T) - \psi(\theta))^2] \\ &= E[(T - E(T))^2 + (E(T) - \psi(\theta))^2 + 2(T - E(T))(E(T) - \psi(\theta))] \\ &= E[(T - E(T))^2] + (E(T) - \psi(\theta))^2 + 2E[T - E(T)](E(T) - \psi(\theta)) \\ &= E[(T - E(T))^2] + (E(T) - \psi(\theta))^2 \\ &\quad \text{(Since } E[T - E(T)] = E(T) - E(T) = 0) \\ &= Var(T) + (E(T) - \psi(\theta))^2 \\ &= Var(T) + Bias^2(T) \end{aligned}$$

■

Bias The bias of an estimator T of $\psi(\theta)$ is given by

$$E_{\theta}(T) - \psi(\theta)$$

Unbiased estimator: When the bias of an estimator is zero, it's called unbiased

Remark

1. For unbiased estimators,

$$MSE_{\theta}(T) = Var_{\theta}(T)$$

2. If all the other properties are similar, then an unbiased estimator is preferred over a biased estimator.
3. In practice, often an biased estimator with lower variance is preferred over an unbiased estimator with really high variance. **We minimize MSE.**

8 Population Variance (σ^2)

Definition $\sigma^2 = E[(X - \mu)^2]$ where $\mu = E[X]$.

If we have equally likely N data points in our population, this is equivalent of

$$\sigma^2 = \frac{1}{N} \sum_{i=1}^N (X_i - \mu)^2$$

In words: It's the average squared difference of each of the data points (X_i) from the mean (μ)

Estimate σ^2 based on a sample of size n When we are estimating based on the sample of size n , we replace μ by \bar{X} , so the numerator is $\sum_{i=1}^n (X_i - \bar{X})^2$. We can divide it by both n or $n - 1$. The latter one is unbiased!

The fraction, $\frac{n-1}{n} \rightarrow 1$ as $n \rightarrow \infty$. So for large n , both estimator will produce similar estimate. In statistical literature, whenever we say *sample variance* we refer to the *unbiased* one. Hence, from now on,

Definition of sample variance

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$$

9 Sampling distribution of S^2 (under Normal Distribution)

Theorem Suppose $X_1, X_2, \dots, X_n \sim N(\mu, \sigma^2)$ iid, $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ and $S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2$.
Then

1. \bar{X} and S^2 are independent, and
2. $\frac{(n-1)S^2}{\sigma^2} \sim \chi^2_{(df=n-1)}$

proof:

Part 1 Let

$$\begin{aligned} U = \bar{X} &= \frac{1}{n}X_1 + \dots + \frac{1}{n}X_n \\ V = X_1 - \bar{X} &= X_1 - \left(\frac{1}{n}X_1 + \dots + \frac{1}{n}X_n\right) \\ &= \left(1 - \frac{1}{n}\right)X_1 - \left(\frac{1}{n}X_2 + \dots + \frac{1}{n}X_n\right) \end{aligned}$$

$$\begin{aligned} Cov(\bar{X}, X_1 - \bar{X}) &= Cov(\bar{X}, X_1) - Cov(\bar{X}, \bar{X}) \\ &= Cov\left(\frac{1}{n}X_1 + \dots + \frac{1}{n}X_n, X_1\right) - \frac{\sigma^2}{n} \\ &= \frac{1}{n}Cov(X_1, X_1) - \frac{\sigma^2}{n} \\ &= \frac{\sigma^2}{n} - \frac{\sigma^2}{n} \\ &= 0 \end{aligned}$$

Hence by E&R theorem, U and V are independent. Similarly, we can show \bar{X} is independent to each $X_i - \bar{X}$ for $i = 1, \dots, n$

Therefore, \bar{X} is independent to $\sum_{i=1}^n (X_i - \bar{X})^2$

Therefore, \bar{X} is independent to $\frac{\sum_{i=1}^n (X_i - \bar{X})^2}{n-1} = S^2$

Part 2

$$\begin{aligned}
 \sum_i (X_i - \mu)^2 &= \sum_i (X_i - \bar{X} + \bar{X} - \mu)^2 \\
 &= \sum_i (X_i - \bar{X})^2 + \sum_i (\bar{X} - \mu)^2 + 2 \sum_i (X_i - \bar{X})(\bar{X} - \mu) \\
 &= \sum_i (X_i - \bar{X})^2 + \sum_i (\bar{X} - \mu)^2 + 2(\bar{X} - \mu) \sum_i (X_i - \bar{X}) \\
 &= \sum_i (X_i - \bar{X})^2 + \sum_i (\bar{X} - \mu)^2 + 2(\bar{X} - \mu) \left(\sum_i X_i - n\bar{X} \right) \\
 &= \sum_i (X_i - \bar{X})^2 + \sum_i (\bar{X} - \mu)^2 + 2(\bar{X} - \mu)(n\bar{X} - n\bar{X}) \\
 &= \sum_i (X_i - \bar{X})^2 + n(\bar{X} - \mu)^2 \\
 \implies \sum_i (X_i - \bar{X})^2 &= \sum_i (X_i - \mu)^2 - n(\bar{X} - \mu)^2
 \end{aligned}$$

$$\begin{aligned}
 \frac{\sum_{i=1}^n (X_i - \mu)^2}{\sigma^2} &= \frac{\sum_i (X_i - \bar{X})^2}{\sigma^2} + \frac{n(\bar{X} - \mu)^2}{\sigma^2} \\
 \implies \sum_i \left(\frac{X_i - \mu}{\sigma} \right)^2 &= \frac{\sum_i (X_i - \bar{X})^2}{\sigma^2} + \left(\frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \right)^2 \\
 &= \frac{(n-1)S^2}{\sigma^2} + \left(\frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \right)^2 \\
 \implies \chi_{(n)}^2 &= \frac{(n-1)S^2}{\sigma^2} + \chi_{(1)}^2 \\
 \implies MGF(\chi_{(n)}^2) &= MGF\left(\frac{(n-1)S^2}{\sigma^2} + \chi_{(1)}^2\right) \\
 &= MGF\left(\frac{(n-1)S^2}{\sigma^2}\right) * MGF(\chi_{(1)}^2) \\
 \implies MGF\left(\frac{(n-1)S^2}{\sigma^2}\right) &= \frac{MGF(\chi_{(n)}^2)}{MGF(\chi_{(1)}^2)} \\
 &= \frac{(1-2t)^{-\frac{n}{2}}}{(1-2t)^{-\frac{1}{2}}} \\
 &= (1-2t)^{-\frac{n-1}{2}}
 \end{aligned}$$

which is the MGF of $\chi_{(n-1)}^2$ ■

E&R theorem 4.6.2 $X_1, X_2, \dots, X_n \sim N(\mu, \sigma^2) i.i.d.$, U and V are two different linear combinations of the X_i 's, then
 $Cov(U, V) = 0 \iff$ U and V are independent.

Note In general, zero covariance doesn't imply independent
 Example: $X \sim N(0, 1)$, $Y = X^2$, clearly X and Y are dependent, but

$$\begin{aligned} Cov(X, Y) &= E[XY] - E[X]E[Y] \\ &= E[X^3] - 0 \cdot E[Y] \\ &= E[X^3] \\ &= \int x^3 f(x) dx \\ &= 0 \quad (\text{since } x^3 f(x) \text{ is centro-symmetric}) \end{aligned}$$

Unbiasedness of S^2 using the Chi-sq distribution

$$\begin{aligned} E\left[\frac{(n-1)S^2}{\sigma^2}\right] &= n-1 \\ \implies E[S^2] &= \sigma^2 \end{aligned}$$

This proves S^2 is an unbiased estimator for σ^2 under Normal distribution
 There's another way to prove it under any arbitrary distribution with the assumption that X_i 's are i.i.d. and μ, σ^2 exists.

10 Some relationships among distributions

1. $\frac{\bar{X} - \mu}{S/\sqrt{n}} \sim t_{(n-1)}$
2. $\frac{\chi_{(m)}^2}{m} \xrightarrow{P} 1$

11 Difference between sample variance and variance of sample mean

variance of sample mean: Expectation of squared difference of sample mean from the true mean

sample variance: average squared difference of each data points in the sample from the sample mean

12 Consistent Estimator

Definition Let T_n be an estimator of parameter θ , T_n is said to be consistent (in probability) if

$$T_n \xrightarrow{P} \theta$$

In words, T_n converges to θ in probability.

Note If $T_n \xrightarrow{a.s.} \theta$ then T_n is called consistent (almost surely). In this course we will only talk about consistent (in probability)

Proving consistency using LLN LLN tells us, $\bar{X} = \frac{1}{n} \sum X_i \xrightarrow{P} E[X_i]$ for any distribution. Immediately that tells us:

1. If $X_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$ then \bar{X} is a consistent estimator of μ
2. If $X_i \stackrel{iid}{\sim} Poisson(\lambda)$ then \bar{X} is a consistent estimator of λ
3. And we can say this for few other known distributions

Goal: prove consistency when the estimator is not simply \bar{X} Still use LLN but with the help of a well known Lemma and the continuous mapping theorem

Slutsky's Lemma We have two different sequence X_n and Y_n

If $X_n \xrightarrow{P} X$ and $Y_n \xrightarrow{P} Y$, then $X_n + Y_n \xrightarrow{P} X + Y$

If $X_n \xrightarrow{P} X$ and $Y_n \xrightarrow{P} Y$, then $X_n Y_n \xrightarrow{P} XY$

Continuous mapping theorem Let $X_n \xrightarrow{P} X$ and $g()$ be a continuous function, then $g(X_n) \xrightarrow{P} g(X)$

Proving S^2 is a consistent estimator of σ^2 ...

MSE consistent An estimator T_n is called MSE consistent if

$$MSE(T_n) \rightarrow 0 \text{ as } n \rightarrow \infty$$

Example: for $N(\mu, \sigma^2)$ $MSE(\bar{X}) = \sigma^2/n \rightarrow 0$ as $n \rightarrow \infty$ Therefore \bar{X} is a MSE consistent estimator of μ

In naive words, after you have calculated the MSE of an estimator, just check if it goes to zero for large n

Note MSE consistent \implies consistent (in probability)

13 Efficient Estimator

Definition of Efficiency Let T_1 and T_2 be two different estimators of θ , Efficiency of T_1 relative to T_2 is defined as

$$eff(T_1, T_2) = \frac{var[T_2]}{var[T_1]}$$

Remark

1. $eff(T_1, T_2) > 1 \implies T_1$ has smaller variance $\implies T_1$ is more efficient
2. This comparison is meaningful when T_1 and T_2 are both unbiased or both have the same bias.

Lower bound of the variance of an unbiased estimator This famous inequality provides a lower bound for the variance of all the unbiased estimators. In other words it gives a lower bound of the MSE (since Bias = 0). The estimator whose variance achieves this lower bound is said to be efficient. Before we state the inequality let's define few terms...

Score function, $S(\theta)$ The derivative of the log-likelihood

$$S(\theta) = \frac{\partial l(\theta)}{\partial \theta}$$

For the random variable X , $S(\theta|X = x) = \frac{\partial}{\partial \theta} \ln f_{\theta}(x)$. For an observed i.i.d sample, it's written as $S(\theta|x_1, x_2, \dots, x_n)$ with

$$S(\theta|x_1, x_2, \dots, x_n) = \frac{\partial}{\partial \theta} \sum_i \ln f_{\theta}(x_i) = \sum_i \frac{\partial}{\partial \theta} \ln f_{\theta}(x_i) = \sum_i S(\theta|x_i)$$

Fisher Information, $I(\theta)$ The function

$$I(\theta) = var_{\theta}[S(\theta|X)]$$

It's the amount of information that each observable random variable X contains about θ .

Information of a sample of size $n = var[S(\theta|x_1, x_2, \dots, x_n)] = nI(\theta)$

A plot showing the randomness of $S(\theta)$ The likelihood function looks different for different data!

One important property of $S(\theta)$ Under some assumptions,

$$E[S(\theta|X = x)] = 0$$

Which implies

$$E[S(\theta|x_1, x_2, \dots, x_n)] = \sum_i E[S(\theta|x_i)] = 0$$

Cramer-Rao Inequality Let X_1, X_2, \dots, X_n be i.i.d. with density $f_\theta(x)$, $T(X_1, X_2, \dots, X_n)$ be an unbiased estimator of θ , Then under some assumptions on $f_\theta(x)$,

$$\text{var}[T] \geq \frac{1}{nI(\theta)}$$

$\frac{1}{nI(\theta)}$ is also known as the Cramer-Rao lower bound (CRLB)

Proof of Cramer-Rao Inequality ...

Definition of sufficient statistic A statistic $T(X_1, X_2, \dots, X_n)$ is said to be sufficient for θ if the conditional distribution of X_1, X_2, \dots, X_n , given $T = t$, does not depend on θ